

- EXL Services — AI Cloud Architect (L1 Interview Preparation)
 - Table of Contents
 - 1. AWS Cloud Architecture (Core)
 - Q1: Describe your experience designing multi-account AWS architectures. How would you structure an enterprise landing zone?
 - Q2: How do you design a highly available, fault-tolerant architecture on AWS?
 - Q3: Explain VPC design for a large enterprise with 200+ microservices.
 - Q4: How do you handle networking between on-premises and AWS?
 - Q5: What is AWS Well-Architected Framework and how do you apply it in practice?
 - 2. Generative AI & LLM Architecture on AWS
 - Q6: How would you architect an enterprise Generative AI solution on AWS?
 - Q7: Explain RAG (Retrieval-Augmented Generation) and how you'd implement it on AWS.
 - Q8: What is Amazon Bedrock and how does it compare to using SageMaker for GenAI?
 - Q9: How do you handle prompt engineering and prompt management in production?
 - Q10: What are Bedrock Guardrails and why are they critical for BFSI?
 - 3. Amazon SageMaker — Deep Dive
 - Q11: Describe your hands-on experience with Amazon SageMaker.
 - Q12: Explain SageMaker Pipelines and how you've used them for MLOps.
 - Q13: How do you handle model training at scale on SageMaker?
 - Q14: How do you deploy SageMaker models to production with zero downtime?
 - Q15: What is SageMaker Feature Store and why is it important?
 - Q16: How do you fine-tune a Large Language Model using SageMaker?
 - Q17: What is SageMaker Model Monitor and how do you detect model drift?
 - 4. AI/ML Pipeline & MLOps
 - Q18: How do you design an end-to-end MLOps pipeline on AWS?
 - Q19: How do you handle data versioning and lineage in ML workflows?
 - 5. Solution Architecture & Design Patterns
 - Q20: How would you architect an AI-powered document processing solution for a bank?

- Q21: Design a real-time fraud detection system using AWS AI services.
- Q22: How would you design a conversational AI platform (chatbot) for an insurance company?
- 6. Security, Compliance & Governance (BFSI)
 - Q23: How do you secure AI/ML workloads on AWS for BFSI clients?
 - Q24: Explain the AWS Shared Responsibility Model in the context of AI/ML.
 - Q25: How do you implement model governance and responsible AI?
- 7. Cost Optimization & FinOps
 - Q26: How do you optimize costs for AI/ML workloads on AWS?
 - Q27: How do you implement FinOps for AWS AI/ML at an organizational level?
- 8. Data Architecture & Analytics
 - Q28: How do you design a data lake on AWS for AI/ML workloads?
 - Q29: How do you handle real-time data streaming for AI/ML on AWS?
- 9. DevOps & Infrastructure as Code
 - Q30: How do you manage AWS infrastructure for AI/ML using Terraform?
 - Q31: How do you implement CI/CD for ML models?
- 10. Behavioral & Leadership (L1 Focus)
 - Q32: Tell me about yourself and your experience as a Cloud/AI Architect.
 - Q33: Describe a complex architecture you designed. What were the trade-offs?
 - Q34: How do you handle a situation where the client wants a quick PoC but you know the architecture won't scale?
 - Q35: How do you stay current with rapidly evolving AI/cloud technologies?
 - Q36: How do you handle disagreements with development teams about architecture decisions?
 - Q37: How do you mentor junior architects and engineers?
- 11. EXL-Specific Questions
 - Q38: Why EXL Services?
 - Q39: What value would you bring as an AI Cloud Architect at EXL?
 - Q40: How would you approach a new BFSI client wanting to adopt GenAI?
- Quick Reference — Top SageMaker Services to Highlight
- Interview Tips for EXL L1 Round

EXL Services — AI Cloud Architect (L1 Interview Preparation)

Role: AI Cloud Architect | **Company:** EXL Services

Experience: 11-15 years | **Location:** Noida/Gurgaon/Pune/Bangalore

Focus Areas: AWS Cloud, Generative AI, SageMaker, BFSI Domain

Prepared by: Pushparaj Naik

Table of Contents


1. [AWS Cloud Architecture \(Core\)](#)
 2. [Generative AI & LLM Architecture on AWS](#)
 3. [Amazon SageMaker — Deep Dive](#)
 4. [AI/ML Pipeline & MLOps](#)
 5. [Solution Architecture & Design Patterns](#)
 6. [Security, Compliance & Governance \(BFSI\)](#)
 7. [Cost Optimization & FinOps](#)
 8. [Data Architecture & Analytics](#)
 9. [DevOps & Infrastructure as Code](#)
 10. [Behavioral & Leadership \(L1 Focus\)](#)
 11. [EXL-Specific Questions](#)
-

1. AWS Cloud Architecture (Core)

Q1: Describe your experience designing multi-account AWS architectures. How would you structure an enterprise landing zone?

Answer:

I've designed multi-account architectures using AWS Control Tower and AWS Organizations. My approach follows the **AWS Well-Architected Framework** with a clear account separation strategy:



Key design decisions:

- **Transit Gateway** for hub-and-spoke networking across accounts
- **AWS SSO (IAM Identity Center)** for centralized authentication with Entra ID/Okta federation
- **Service Control Policies (SCPs)** to enforce guardrails (e.g., deny root user, restrict regions)
- **AWS Config** with conformance packs for continuous compliance
- **Terraform** with remote state in S3 + DynamoDB locking for IaC

For an EXL client in BFSI, I'd add dedicated accounts for PCI-DSS workloads and a separate AI/ML account with SageMaker to isolate model training costs and data access.

Q2: How do you design a highly available, fault-tolerant architecture on AWS?

Answer:

High availability on AWS is about **eliminating single points of failure** across every layer:

Layer	HA Pattern	AWS Service
Compute	Multi-AZ Auto Scaling Groups, EKS with node groups across 3 AZs	EC2 ASG, EKS, ECS

Layer	HA Pattern	AWS Service
Database	Multi-AZ RDS with automatic failover, Aurora Global Database for cross-region	RDS, Aurora, DynamoDB Global Tables
Caching	ElastiCache Redis with Multi-AZ and automatic failover	ElastiCache
Storage	S3 (11 9's durability), cross-region replication for DR	S3, EFS
Networking	Multi-AZ ALB/NLB, Route 53 health checks with failover routing	ALB, Route 53
Messaging	SQS (distributed), SNS fan-out, EventBridge	SQS, SNS

Real example: For a BFSI client, I architected a payment processing system with:

- Aurora PostgreSQL Multi-AZ (RPO=0, RTO < 30s)
- Active-passive DR with Aurora Global Database (us-east-1 → eu-west-1)
- DynamoDB Global Tables for session state (multi-region active-active)
- Route 53 with health-check-based failover
- Chaos engineering with AWS Fault Injection Simulator for DR drills

Q3: Explain VPC design for a large enterprise with 200+ microservices.

Answer:

For a large microservices estate, I use a **hub-and-spoke VPC topology** with Transit Gateway:

Hub VPC (Shared Services):

- Transit Gateway attachment
- Centralized NAT Gateways
- VPN/Direct Connect termination
- Centralized DNS (Route 53 Resolver)

CIDR: 10.0.0.0/16

Spoke VPCs (per environment/team):

- Production VPC: 10.1.0.0/16 (3 AZs × 3 tiers)
- Staging VPC: 10.2.0.0/16
- Dev VPC: 10.3.0.0/16

– AI/ML VPC: 10.4.0.0/16 (SageMaker endpoints)

Subnet Strategy (per VPC):

- Public (10.x.1.0/24, 10.x.2.0/24, 10.x.3.0/24) – ALB, NAT GW
- Private (10.x.11.0/24, 10.x.12.0/24, 10.x.13.0/24) – EKS, EC2
- Data (10.x.21.0/24, 10.x.22.0/24, 10.x.23.0/24) – RDS, ElastiCache

Security controls:

- **Security Groups:** Service-level (e.g., `sg-storefront-api` → `sg-aurora-primary` on port 5432)
- **NACLs:** Subnet-level deny rules for defense-in-depth
- **VPC Endpoints:** Gateway (S3, DynamoDB) + Interface (SageMaker, Secrets Manager, ECR) to avoid internet traversal
- **PrivateLink:** For cross-account SageMaker endpoint access
- **Flow Logs:** Encrypted with KMS, shipped to CloudWatch + S3 for audit

Q4: How do you handle networking between on-premises and AWS?

Answer:

I've implemented hybrid connectivity using multiple patterns:

1. **AWS Direct Connect** (dedicated 10Gbps): Primary path for production workloads with low-latency requirements
2. **Site-to-Site VPN** (IPSec over internet): Backup path with BGP routing for automatic failover
3. **Transit Gateway:** Central hub connecting VPCs + Direct Connect + VPN



For DNS: Route 53 Resolver inbound/outbound endpoints for bi-directional name resolution between on-prem and AWS.

Q5: What is AWS Well-Architected Framework and how do you apply it in practice?

Answer:

The Well-Architected Framework has **6 pillars**:

Pillar	My Practical Application
Operational Excellence	IaC (Terraform), CI/CD (GitHub Actions), runbooks, CloudWatch dashboards
Security	Zero-trust networking, KMS CMK encryption, IAM least-privilege, GuardDuty
Reliability	Multi-AZ deployments, auto-scaling, chaos engineering, DR drills
Performance Efficiency	Right-sizing (Compute Optimizer), caching (ElastiCache), async processing (SQS)
Cost Optimization	Reserved Instances, Savings Plans, Spot for SageMaker training, S3 lifecycle policies
Sustainability	Graviton instances (ARM), serverless where possible, right-sized resources

I conduct Well-Architected Reviews quarterly using the AWS Well-Architected Tool and track remediation of high/medium findings. For EXL's BFSI clients, the Security and Reliability pillars receive the most focus.

2. Generative AI & LLM Architecture on AWS

Q6: How would you architect an enterprise Generative AI solution on AWS?

Answer:

I've designed GenAI architectures using **Amazon Bedrock** as the foundation with SageMaker for custom model fine-tuning:

ENTERPRISE GENAI ARCHITECTURE ON AWS:

APPLICATION LAYER

- └ API Gateway (REST/WebSocket for streaming)
- └ Lambda / ECS (orchestration layer)
- └ CloudFront (static assets, response caching)

ORCHESTRATION LAYER

- └ LangChain / LlamaIndex (RAG pipeline)
- └ Amazon Bedrock Agents (tool use, multi-step reasoning)
- └ Step Functions (complex workflow orchestration)

MODEL LAYER

- └ Amazon Bedrock (Claude 3.5, Titan, Llama 3)
- └ SageMaker Endpoints (custom fine-tuned models)
- └ SageMaker JumpStart (pre-trained model catalog)

DATA / KNOWLEDGE LAYER

- └ Amazon Bedrock Knowledge Bases (managed RAG)
- └ OpenSearch Serverless (vector store, k-NN)
- └ S3 (document corpus – PDFs, policies, FAQs)
- └ RDS PostgreSQL + pgvector (transactional + embeddings)

GUARDRAILS & GOVERNANCE

- └ Amazon Bedrock Guardrails (content filters, PII blocking)
- └ CloudWatch (latency, token usage, cost tracking)
- └ SageMaker Model Monitor (data drift, quality)
- └ AWS CloudTrail (audit logging for all model invocations)

For a BFSI use case at EXL, I'd add:

- **Bedrock Guardrails** to block PII/financial data in prompts
 - **VPC endpoints** for Bedrock to ensure no internet traversal
 - **Prompt injection detection** middleware
 - **Human-in-the-loop** approval for any customer-facing responses
-

Q7: Explain RAG (Retrieval-Augmented Generation) and how you'd implement it on AWS.

Answer:

RAG enhances LLM responses by retrieving relevant context from a knowledge base before generation, reducing hallucinations and providing up-to-date information.

My AWS RAG Implementation:

DOCUMENT INGESTION PIPELINE:

S3 (docs) → Lambda (chunking) → Bedrock Embeddings (Titan) → OpenSearch Serverless (vector store)

QUERY FLOW:

User Query → Embedding → Vector Search (top-k) → Augmented Prompt → LLM → Response

Step-by-step:

1. INGEST: PDF/Word docs uploaded to S3 → triggers Lambda
2. CHUNK: Text split into 512-token chunks with 50-token overlap
3. EMBED: Amazon Titan Embedding V2 (1024 dimensions)
4. STORE: OpenSearch Serverless with k-NN index (HNSW algorithm)
5. QUERY: User question → embed → cosine similarity search (top 5 chunks)
6. AUGMENT: System prompt + retrieved chunks + user question
7. GENERATE: Claude 3.5 Sonnet via Bedrock → streaming response
8. CITE: Include source document references in response

Alternatively, using managed RAG:

- **Amazon Bedrock Knowledge Bases** handles steps 1-5 automatically
- I connect S3 as the data source, select embedding model, and Bedrock manages chunking, embedding, and vector storage in OpenSearch Serverless
- This reduces operational overhead significantly

For BFSI: RAG is ideal for policy Q&A, regulatory compliance queries, and customer support chatbots where accuracy is critical and hallucination is unacceptable.

Q8: What is Amazon Bedrock and how does it compare to using SageMaker for GenAI?

Answer:

Aspect	Amazon Bedrock	SageMaker for GenAI
What it is	Managed API to access foundation models (Claude, Titan, Llama, Mistral)	Full ML platform for training, fine-tuning, and deploying custom models
Use when	You need out-of-the-box LLM capabilities, RAG, or agents	You need custom model training, fine-tuning on proprietary data, or full control
Fine-tuning	Continued pre-training & fine-tuning (limited models)	Full fine-tuning with LoRA, QLoRA, PEFT on any model using custom scripts
Infrastructure	Fully managed (no instances to manage)	You manage instance types (ml.p4d, ml.p5, etc.)
Cost model	Pay-per-token (input/output) or Provisioned Throughput	Pay for training hours + endpoint instance hours
Governance	Bedrock Guardrails (built-in PII/content filtering)	SageMaker Model Monitor + custom guardrails
Best for EXL	Client-facing chatbots, document Q&A, summarization	Custom fraud detection models, industry-specific fine-tuned LLMs

My approach: I use **Bedrock for inference** (quick time-to-value) and **SageMaker for training/fine-tuning** when clients need custom models trained on their proprietary data. They complement each other — SageMaker trains the model, Bedrock hosts the custom model for inference at scale.

Q9: How do you handle prompt engineering and prompt management in production?


Answer:

Prompt engineering in production requires systematic management:

1. **Prompt Templates:** Version-controlled in Git (YAML/JSON), not hardcoded

- 2. **Prompt Registry:** Central store (DynamoDB or Parameter Store) for A/B testing different prompts
- 3. **System Prompts:** Define persona, guardrails, output format in system-level instructions
- 4. **Few-shot Examples:** Include 2-3 examples in prompts for consistent output formatting
- 5. **Chain-of-Thought:** For complex reasoning, instruct the model to "think step by step"

Production prompt management pipeline:



BFSI-specific guardrails I embed in prompts:

- "Never provide specific financial advice or investment recommendations"
- "Always include a disclaimer when discussing insurance policies"
- "If the question is about a specific account, respond with 'Please contact your account manager'"
- "Never reveal internal process details or system architecture"

Q10: What are Bedrock Guardrails and why are they critical for BFSI?

Answer:

Bedrock Guardrails provide configurable safety filters for LLM applications:

Guardrail Type	Purpose	BFSI Relevance
Content Filters	Block harmful/inappropriate content by category (hate, violence, sexual, etc.)	Compliance with financial regulations

Guardrail Type	Purpose	BFSI Relevance
Denied Topics	Block specific topics (e.g., "Give me insider trading advice")	Prevent regulatory violations
Word Filters	Block specific words or phrases	Block competitor names, internal jargon
PII Filters	Detect and mask PII (SSN, credit card, account numbers)	GDPR, PCI-DSS, SOX compliance
Contextual Grounding	Check if response is grounded in provided context (anti-hallucination)	Accuracy in policy/regulatory Q&A

For BFSI at EXL, I configure:

- **PII masking:** Automatically redact SSN, account numbers, credit card numbers from both input and output
- **Denied topics:** Investment advice, credit decisions, insurance underwriting
- **Grounding checks:** Ensure all responses cite source documents (RAG grounding)
- **CloudTrail logging:** Every guardrail evaluation is audit-logged for compliance

3. Amazon SageMaker — Deep Dive

Q11: Describe your hands-on experience with Amazon SageMaker.

Answer:

I've worked extensively with SageMaker across the ML lifecycle:

Training:

- Set up distributed training jobs using **SageMaker Training** with **ml.p4d.24xlarge** instances for transformer fine-tuning
- Used **SageMaker Experiments** to track hyperparameters, metrics, and artifacts across 100+ training runs
- Implemented **Spot Training** (managed spot instances) — reduced training costs by 60-70%

- Used **SageMaker Debugger** to detect vanishing gradients and overfitting during training

Model Development:

- **SageMaker Studio** as the primary IDE for our data science team
- **SageMaker JumpStart** to evaluate pre-trained foundation models (BERT, GPT-J, Falcon) before fine-tuning
- **SageMaker Processing** for feature engineering (Spark/SKLearn containers)
- **SageMaker Feature Store** (offline + online) for centralized feature management

Deployment:

- **Real-time Endpoints** for fraud detection (< 100ms latency requirement)
- **Serverless Inference** for infrequent batch scoring (cost optimization)
- **Multi-Model Endpoints** to host 50+ customer-specific models on single infrastructure
- **Asynchronous Inference** for document processing (large payload, long running)

Monitoring:

- **SageMaker Model Monitor** for data drift detection (population stability index)
- **Model quality monitoring** — tracking accuracy/F1 degradation over time
- **CloudWatch custom metrics** for business KPIs (fraud detection rate, false positive rate)


Q12: Explain SageMaker Pipelines and how you've used them for MLOps.

Answer:

SageMaker Pipelines is a CI/CD system purpose-built for ML workflows. I've used it to automate end-to-end model lifecycle:

MY SAGEMAKER PIPELINE STRUCTURE:





Pipeline steps I implemented:

1. **QualityCheck:** Data Wrangler checks for schema drift, missing values, outliers
2. **Processing:** SageMaker Processing job (Spark) for feature engineering from S3 data lake
3. **Training:** XGBoost/Hugging Face training with hyperparameter tuning (HPO)
4. **Evaluation:** Model evaluation against test set — output metrics to JSON
5. **Condition:** If $F1 > \text{threshold}$ → proceed to registration, else → stop + alert
6. **RegisterModel:** Push to SageMaker Model Registry with metadata (accuracy, data version, owner)
7. **Deploy:** Create/update SageMaker endpoint (blue/green or canary)

Integration with GitHub Actions:

```
Git push → GitHub Actions → trigger SageMaker Pipeline → Model Registry → approval → deploy
```

I used **SageMaker Model Registry** with approval workflow — Data Scientists approve model in "PendingManualApproval" status before it's promoted to production.

Q13: How do you handle model training at scale on SageMaker?

Answer:

For large-scale training, I use several SageMaker capabilities:

1. Distributed Training:

- **Data Parallelism:** SageMaker's distributed data parallel library — splits mini-batches across GPUs
- **Model Parallelism:** For models that don't fit in single GPU memory — pipeline and tensor parallelism
- Instance: `ml.p4d.24xlarge` (8 × A100 GPUs, 40GB each) for transformer training

2. Cost Optimization:

- **Managed Spot Training:** Use spot instances with checkpointing to S3 — 60-70% cost reduction
- **SageMaker Savings Plans:** 1-year commitments for predictable training workloads
- **Right-sizing:** Start with `ml.m5.xlarge` for prototyping, scale to GPU instances for final training

3. Hyperparameter Tuning:

- **SageMaker Automatic Model Tuning (HPO):**
 - Bayesian optimization strategy for efficient exploration
 - Run 50+ concurrent trials with early stopping
 - Objective metric: maximize F1-score or minimize validation loss

4. Training with Custom Containers:

```
# SageMaker Training with Hugging Face for LLM fine-tuning
from sagemaker.huggingface import HuggingFace

huggingface_estimator = HuggingFace(
    entry_point='train.py',
    source_dir='./scripts',
    instance_type='ml.p4d.24xlarge',
    instance_count=2,
    role=sagemaker_role,
    transformers_version='4.37',
    pytorch_version='2.1',
    py_version='py310',
    hyperparameters={
        'model_name': 'meta-llama/Llama-2-7b-hf',
        'epochs': 3,
        'per_device_train_batch_size': 4,
        'gradient_accumulation_steps': 8,
        'learning_rate': 2e-5,
```

```

        'lora_r': 16,                # LoRA rank
        'lora_alpha': 32,
    },
    use_spot_instances=True,
    max_wait=72000,
    checkpoint_s3_uri=f's3://{bucket}/checkpoints/',
)

huggingface_estimator.fit({'train': train_data, 'test': test_data})

```

Q14: How do you deploy SageMaker models to production with zero downtime?

Answer:

I use multiple SageMaker deployment strategies based on the use case:

1. Blue/Green Deployment (Default — Zero Downtime):

```

from sagemaker.model import Model

model = Model(
    model_data=model_uri,
    role=role,
    image_uri=inference_image,
)

# Deploy with blue/green (automatic with UpdateEndpoint)
predictor = model.deploy(
    initial_instance_count=2,
    instance_type='ml.m5.xlarge',
    endpoint_name='fraud-detection-prod',
    update_endpoint=True, # Blue/green behind the scenes
)

```

2. Canary Deployment (Shadow Testing):

- Route 10% traffic to new model variant, 90% to existing
- Compare metrics (latency, accuracy) before full switchover
- SageMaker Production Variants support this natively

3. Multi-Model Endpoints (Cost Efficiency):

- Host 50+ customer-specific models on shared infrastructure
- Models loaded/unloaded from S3 on demand

- Ideal for multi-tenant BFSI scenarios (each bank client has own model)

4. Serverless Inference (Infrequent Workloads):

- No instances running when idle — scales to zero
- Cold start ~30s, but perfect for batch scoring APIs used a few times/day

Monitoring post-deploy:

- SageMaker Model Monitor (data drift + model quality)
 - CloudWatch alarms on **Invocation4XXErrors**, **ModelLatency**, **OverheadLatency**
 - Custom A/B metrics via CloudWatch custom metrics
-

Q15: What is SageMaker Feature Store and why is it important?

Answer:

SageMaker Feature Store is a centralized repository for ML features, providing consistency between training and inference:

Component	Purpose	My Usage
Offline Store	S3-backed, for batch training data	Feature engineering pipelines write Parquet to offline store
Online Store	Low-latency (< 10ms) key-value lookups	Real-time fraud detection reads customer features at inference
Feature Groups	Logical grouping of related features	customer-features , transaction-features , device-features

Why it matters:

- **Training-serving skew prevention:** Same feature definitions used in training and inference
- **Feature reuse:** 30+ features computed once, shared across fraud, credit, and churn models
- **Time-travel:** Query features as they existed at a specific point in time (critical for backtesting)

- **Data lineage:** Know which features were used to train which model version

Real example: For a BFSI fraud detection pipeline:

Online Feature Store (real-time):

- customer_id → avg_transaction_30d, total_orders_90d, device_fingerprint_count
- Low-latency lookup at inference time (< 10ms via GetRecord API)

Offline Feature Store (training):

- S3/Parquet with historical feature snapshots
- Used by SageMaker Training job for model retraining (weekly)

Q16: How do you fine-tune a Large Language Model using SageMaker?

Answer:

I've fine-tuned LLMs on SageMaker using several approaches:

1. SageMaker JumpStart (Easiest):

- Select a base model (Llama 2, Falcon, Mistral) from JumpStart
- Provide training data in instruction-following format (JSON Lines)
- JumpStart handles container, distributed training, and deployment
- Good for quick PoCs

2. Hugging Face on SageMaker (My Primary Approach):

```
# Fine-tune Llama 2 with QLoRA (quantized LoRA) for cost efficiency
hyperparameters = {
    'model_id': 'meta-llama/Llama-2-7b-hf',
    'epochs': 3,
    'per_device_train_batch_size': 4,
    'gradient_accumulation_steps': 8,
    'lora_r': 16,
    'lora_alpha': 32,
    'lora_dropout': 0.05,
    'quantization': '4bit', # QLoRA – fit 7B model on single GPU
    'learning_rate': 2e-5,
    'max_seq_length': 2048,
}

# Uses ml.g5.2xlarge (single A10G GPU, 24GB) – much cheaper than p4d
huggingface_estimator = HuggingFace(
```

```
entry_point='finetune_qlora.py',
instance_type='ml.g5.2xlarge',
instance_count=1,
...
)
```

3. Training Data Format:

```
{"instruction": "Summarize this insurance claim", "input": "On 15 March...",
"output": "Claim for water damage..."}
{"instruction": "Extract key entities from this loan application", "input":
"...", "output": "..."}

```

4. Evaluation:

- ROUGE scores for summarization
- Exact match / F1 for extraction tasks
- Human evaluation (domain experts rate 200 samples on 1-5 scale)

Cost optimization: QLoRA fine-tuning on **ml.g5.2xlarge** costs ~
2/hour vs. full fine-tuning on 'ml.p4d.24xlarge' at 37/hour — 95% cost reduction
with minimal quality loss.

Q17: What is SageMaker Model Monitor and how do you detect model drift?

Answer:

SageMaker Model Monitor continuously monitors deployed models for quality degradation:

Four types of monitoring I configure:

Monitor Type	What It Detects	My Configuration
Data Quality	Feature drift — statistical distribution changes	Baseline from training data, compare hourly
Model Quality	Accuracy/F1 degradation	Ground truth labels joined after 24hr delay

Monitor Type	What It Detects	My Configuration
Bias Drift	Fairness metric changes	DPPL, DI metrics on protected attributes
Feature Attribution	SHAP value changes	Feature importance drift from baseline

My monitoring pipeline:

```

SageMaker Endpoint
  ↓ (captures inputs/outputs to S3)
Model Monitor Schedule (hourly)
  ↓
Compare against baseline statistics
  ↓
If violation detected:
  ↓
CloudWatch Alarm → SNS → PagerDuty
  ↓
Trigger SageMaker Pipeline (automated retraining)

```

BFSI example: For fraud detection, I monitor:

- **Data drift:** If **avg_transaction_amount** distribution shifts > 2 standard deviations → alert
- **Model quality:** If precision drops below 0.90 → trigger retraining pipeline
- **Concept drift:** If fraud patterns change (new attack vectors), the model's predictions degrade even with correct features

4. AI/ML Pipeline & MLOps

Q18: How do you design an end-to-end MLOps pipeline on AWS?

Answer:

END-TO-END MLOPS ON AWS:

Data Layer:

```
S3 Data Lake → Glue Catalog → Athena (ad-hoc queries)
↓
Feature Engineering:
  SageMaker Processing (Spark) → Feature Store (online + offline)
↓
Model Development:
  SageMaker Studio → Experiments (tracking) → Notebooks
↓
Training Pipeline:
  SageMaker Pipelines → HP0 → Training (Spot) → Evaluation
↓
Model Registry:
  SageMaker Model Registry → Manual Approval → Production
↓
Deployment:
  SageMaker Endpoint (real-time) / Batch Transform / Serverless
↓
Monitoring:
  Model Monitor → CloudWatch → Auto-Retrain (if drift detected)
↓
CI/CD:
  GitHub → GitHub Actions → SageMaker Pipeline trigger → Deploy
```

Key principles I follow:

1. **Reproducibility:** Every training job is tracked (data version, code version, hyperparams, metrics)
2. **Automation:** Model retraining triggered by data drift or schedule (weekly)
3. **Governance:** Model Registry with approval gates, audit trail, lineage
4. **Testing:** Shadow deployment (canary) before production, A/B testing
5. **Rollback:** One-click rollback to previous model version

Q19: How do you handle data versioning and lineage in ML workflows?

Answer:

Data lineage is critical for regulatory compliance in BFSI (model explainability requirements):

Tool	Purpose	My Usage
S3 Versioning	Track data file versions	Training datasets versioned automatically

Tool	Purpose	My Usage
AWS Glue Data Catalog	Schema registry and metadata	Feature definitions, data types, ownership
SageMaker Experiments	Link data → training → model	Each experiment run records: data S3 URI, Git SHA, metrics
SageMaker Model Registry	Model version → data version mapping	Model v2.1 trained on dataset-v5, code commit abc123
SageMaker Lineage	End-to-end artifact tracking	Visualize: data → processing → training → model → endpoint

For BFSI compliance, regulators may ask: "Show me exactly which data was used to train the credit scoring model v3.2, when, and by whom." SageMaker Lineage provides this complete audit trail.

5. Solution Architecture & Design Patterns


Q20: How would you architect an AI-powered document processing solution for a bank?

Answer:

This is a common BFSI use case — Intelligent Document Processing (IDP):

DOCUMENT PROCESSING PIPELINE:

S3 Upload → EventBridge → Step Functions Orchestration



↓
Bedrock / SageMaker
(Classification +
Summarization)
↓
DynamoDB (structured data)
+ S3 (enriched documents)
↓
Human Review (A2I)
(low-confidence results)

AWS Services used:

- **Amazon Textract:** OCR, table extraction, form parsing (bank statements, loan applications)
- **Amazon Comprehend:** NER (names, dates, policy numbers), sentiment analysis
- **Amazon Bedrock:** Document summarization, Q&A over extracted content
- **Amazon A2I (Augmented AI):** Human-in-the-loop for low-confidence extractions
- **Step Functions:** Orchestrate the multi-step pipeline with error handling
- **DynamoDB:** Store structured extraction results with TTL for retention compliance

Q21: Design a real-time fraud detection system using AWS AI services.

Answer:

REAL-TIME FRAUD DETECTION:

Transaction Event

↓
Amazon Kinesis Data Streams (real-time ingestion)
↓

Lambda / ECS (Feature Assembly)


- Query Feature Store (online)
 - avg_txn_30d, device_cnt
- Query ElastiCache (rules)
 - velocity checks, limits
- Assemble feature vector

↓

SageMaker Endpoint (Real-time)

- Model: XGBoost / Neural Net
- Latency: < 50ms (p99)
- Multi-Model Endpoint (per-bank)

Response: fraud_score (0-1)



Key design decisions:

- **< 100ms end-to-end latency** (Kinesis → Feature Store → SageMaker → response)
- **SageMaker Feature Store Online** for pre-computed customer features (< 10ms lookup)
- **Multi-Model Endpoint** to serve bank-specific models on shared infrastructure
- **SageMaker Model Monitor** for drift detection (fraud patterns evolve quickly)
- **A/B testing** new model versions with SageMaker Production Variants

Q22: How would you design a conversational AI platform (chatbot) for an insurance company?


Answer:

CONVERSATIONAL AI PLATFORM:

Customer Channels:

Web Chat / Mobile App / WhatsApp → Amazon Lex (NLU) + Bedrock (GenAI)

Architecture:



Guardrails:

- └ Bedrock Guardrails (PII masking, denied topics)
- └ Lex confidence threshold ($< 0.7 \rightarrow$ escalate to human)
- └ Audit logging (all conversations $\rightarrow$ S3 $\rightarrow$ Athena)

BFSI-specific features:

- **Policy Q&A:** RAG over insurance policy documents (PDFs in S3)
- **Claims Status:** Integration with backend claims API via Lambda
- **Human Escalation:** When confidence $< 70\%$ or customer requests it
- **Sentiment Detection:** Escalate angry customers to senior agents
- **Compliance:** All conversations logged and auditable (insurance regulations)

6. Security, Compliance & Governance (BFSI)

Q23: How do you secure AI/ML workloads on AWS for BFSI clients?

Answer:

Security Layer	Implementation
Network	VPC endpoints for SageMaker, Bedrock, S3 — no internet traversal
Encryption at rest	KMS CMK for S3, EBS, SageMaker volumes, model artifacts
Encryption in transit	TLS 1.3 everywhere, VPC endpoint policies
Access Control	IAM least-privilege, SageMaker execution roles, Bedrock model access policies
Data Protection	S3 bucket policies, VPC endpoint policies, Macie for PII scanning

Security Layer	Implementation
Model Governance	Model Registry approval workflows, model cards, audit trail
Audit Logging	CloudTrail for all API calls, SageMaker model invocation logging
Compliance	Security Hub (PCI-DSS, SOC2), GuardDuty, Config rules
Prompt Security	Bedrock Guardrails, input sanitization, prompt injection detection

Specific BFSI controls:

1. SageMaker Studio in VPC-only mode (no internet access)
 2. KMS customer-managed keys for all data at rest
 3. SageMaker inter-container encryption for distributed training
 4. S3 access points with VPC endpoint policies (data doesn't leave VPC)
 5. CloudTrail + CloudWatch Logs → SIEM integration (Splunk/QRadar)
 6. Quarterly penetration testing of ML API endpoints

Q24: Explain the AWS Shared Responsibility Model in the context of AI/ML.

Answer:

Responsibility	AWS Manages	You Manage
Infrastructure	Physical security, hardware, hypervisor	VPC configuration, security groups, NACLs
SageMaker Platform	Underlying compute, container runtime, patching	Execution roles, VPC settings, encryption config
Bedrock	Model hosting, infrastructure, availability	Guardrails, prompt security, data privacy
Data	S3 durability, storage infrastructure	Classification, encryption keys, access policies, retention
Models	Infrastructure for training/serving	Model quality, fairness, bias monitoring, accuracy

Responsibility	AWS Manages	You Manage
Compliance	SOC/ISO/PCI-DSS certifications of infrastructure	Application-level compliance, audit trails, data governance

Critical point for BFSI: AWS provides the compliant *infrastructure*, but **you are responsible for what you put on it**. Using PCI-DSS compliant S3 doesn't make your application PCI-compliant — you must implement access controls, encryption, logging, and data handling correctly.

Q25: How do you implement model governance and responsible AI?

Answer:

MODEL GOVERNANCE FRAMEWORK:

1. MODEL CARDS (SageMaker):

- Model purpose and intended use
- Training data description
- Performance metrics by demographic
- Limitations and failure modes
- Approval chain and stakeholders

2. BIAS DETECTION (SageMaker Clarify):

- Pre-training: Class Imbalance (CI), Difference in Proportions (DPL)
- Post-training: Disparate Impact (DI), Equal Opportunity Difference
- Runtime: Continuous bias monitoring

3. EXPLAINABILITY (SageMaker Clarify):

- SHAP values for feature attribution
- Partial dependence plots
- Counterfactual explanations ("if X were different...")

4. APPROVAL WORKFLOW:

Data Scientist → Model Registry → Peer Review → Business Approval → Deploy

For BFSI: Regulators (RBI, OCC, FCA) increasingly require **model explainability** — why did the model reject this loan? SageMaker Clarify provides the tooling, but the governance framework (model risk management, validation, and documentation) must be established organizationally.

7. Cost Optimization & FinOps

Q26: How do you optimize costs for AI/ML workloads on AWS?

Answer:

Strategy	Savings	Application
Spot Instances for Training	60-70%	SageMaker managed spot training with checkpointing
SageMaker Savings Plans	Up to 64%	1-year commitment for predictable inference workloads
Right-sizing	20-40%	Compute Optimizer recommendations, start small and scale
Serverless Inference	80%+	For infrequent model invocations (batch scoring APIs)
Multi-Model Endpoints	60-80%	Host 50+ models on shared GPU infrastructure
S3 Intelligent Tiering	30-40%	Training data accessed infrequently moves to IA/Glacier
Instance selection	30-50%	Graviton (ml.c7g) for CPU inference, ml.inf2 for transformer inference

Specific example:

- Moved fraud detection from `ml.m5.xlarge` (0.23/hr) to `ml.inf2.xlarge` (0.76/hr) — 3x throughput, net 50% cost reduction per inference
- Switched LLM training from `ml.p4d.24xlarge` (37/hr) to spot instances — reduced monthly training budget from 15K to \$5K

Q27: How do you implement FinOps for AWS AI/ML at an organizational level?

Answer:

FINOPS FRAMEWORK:

1. VISIBILITY:

- └─ Cost Explorer: Daily spend by SageMaker, Bedrock, S3
- └─ AWS Budgets: Alerts at 80%/100% thresholds
- └─ Cost Allocation Tags: project, team, model-name, environment
- └─ Custom dashboards: Cost per inference, cost per training run

2. OPTIMIZATION:

- └─ Reserved capacity analysis (Compute Optimizer)
- └─ Spot training adoption tracking
- └─ Idle endpoint detection (< 10 invocations/day → consider serverless)
- └─ S3 lifecycle policies (training artifacts → IA after 30 days)

3. GOVERNANCE:

- └─ SCPs: Max instance type limits per account
- └─ IAM: Budget-linked policies (deny CreateEndpoint if budget exceeded)
- └─ Quotas: SageMaker service quotas per team
- └─ Chargeback: Monthly reports to business units

8. Data Architecture & Analytics

Q28: How do you design a data lake on AWS for AI/ML workloads?

Answer:

AWS DATA LAKE ARCHITECTURE:

INGESTION

- └─ Kinesis Data Streams (real-time events)
- └─ Kinesis Firehose (streaming → S3)
- └─ AWS DMS (database replication)
- └─ S3 Transfer Acceleration (file uploads)

STORAGE (S3 – Medallion Architecture)

- └─ Bronze: s3://datalake/raw/ (raw ingestion)
- └─ Silver: s3://datalake/processed/ (cleaned, schema)
- └─ Gold: s3://datalake/curated/ (business-ready)

PROCESSING

- └─ AWS Glue (ETL, schema discovery, catalog)

- └─ EMR (Spark for large-scale transformations)
- └─ SageMaker Processing (ML feature engineering)

ANALYTICS

- └─ Athena (ad-hoc SQL queries on S3)
- └─ Redshift (data warehouse for BI)
- └─ QuickSight (dashboards and visualizations)

GOVERNANCE

- └─ Lake Formation (fine-grained access control)
- └─ Glue Data Catalog (metadata, schema registry)
- └─ Macie (PII detection and classification)

Q29: How do you handle real-time data streaming for AI/ML on AWS?

Answer:

For real-time AI workloads (fraud detection, recommendations):

Component	Service	Purpose
Ingestion	Kinesis Data Streams	Receive events (1M+ events/sec)
Processing	Kinesis Data Analytics (Flink)	Windowed aggregations, feature computation
Feature Store	SageMaker Feature Store (Online)	Low-latency feature serving (< 10ms)
Inference	SageMaker Endpoint (Real-time)	Model scoring (< 100ms)
Actions	Lambda + SNS/SQS	Routing decisions (approve/block/review)
Storage	Kinesis Firehose → S3	Raw event archival for retraining

9. DevOps & Infrastructure as Code

Q30: How do you manage AWS infrastructure for AI/ML using Terraform?

Answer:

I use Terraform to provision all SageMaker and AI infrastructure:

```
# SageMaker Domain (Studio)
resource "aws_sagemaker_domain" "ml_platform" {
  domain_name = "medcloud-ml-platform"
  auth_mode   = "IAM"
  vpc_id      = module.vpc.vpc_id
  subnet_ids  = module.vpc.private_subnet_ids

  default_user_settings {
    execution_role = aws_iam_role.sagemaker_execution.arn

    security_groups = [aws_security_group.sagemaker.id]

    canvas_app_settings {
      direct_deploy_settings {
        status = "ENABLED"
      }
    }
  }
}

app_network_access_type = "VpcOnly" # No internet access
}

# SageMaker Model Endpoint
resource "aws_sagemaker_endpoint_configuration" "fraud_detection" {
  name = "fraud-detection-prod"

  production_variants {
    variant_name      = "primary"
    model_name        = aws_sagemaker_model.fraud_xgboost.name
    instance_type     = "ml.m5.xlarge"
    initial_instance_count = 2
    initial_variant_weight = 1.0
  }

  data_capture_config {
    enable_capture           = true
    initial_sampling_percentage = 100
    destination_s3_uri      =
"s3://${aws_s3_bucket.model_monitor.id}/capture"

    capture_options {
      capture_mode = "Input"
    }
    capture_options {
      capture_mode = "Output"
    }
  }
}
```

```
}  
}
```

Module structure:

```
terraform/  
├── modules/  
│   ├── sagemaker-domain/      # Studio domain + user profiles  
│   ├── sagemaker-endpoint/    # Model serving infrastructure  
│   ├── bedrock-guardrails/    # Guardrail configuration  
│   └── ml-networking/         # VPC endpoints for ML services  
└── environments/  
    ├── dev/      (ml.m5.large, single instance)  
    └── prod/     (ml.m5.xlarge, 2 instances, auto-scaling)
```

Q31: How do you implement CI/CD for ML models?

Answer:

ML CI/CD PIPELINE (GitHub Actions + SageMaker):

1. Data Scientist pushes code to feature branch
2. GitHub Actions triggers:
 - ├── Lint Python code (ruff)
 - ├── Unit tests (pytest)
 - ├── SageMaker Pipeline execution (training + evaluation)
 - └── Model metrics reported back to PR
3. PR Review:
 - ├── Code review by ML Engineer
 - ├── Model metrics review (F1, precision, recall)
 - └── Bias report review (SageMaker Clarify)
4. Merge to main:
 - ├── Model registered in SageMaker Model Registry
 - ├── Approval required (manual gate)
 - └── On approval → SageMaker Endpoint updated (blue/green)
5. Post-deploy:
 - ├── Smoke tests against endpoint
 - ├── Model Monitor enabled
 - └── CloudWatch dashboard updated

10. Behavioral & Leadership (L1 Focus)

Q32: Tell me about yourself and your experience as a Cloud/AI Architect.

Answer:

"I'm a Cloud and AI Architect with 12+ years of experience designing and building enterprise-grade solutions on AWS. My career spans infrastructure automation, cloud migration, and most recently, AI/ML platform engineering.

Key highlights:

- **AWS Architecture:** Designed multi-account landing zones, VPC networking, EKS-based microservices platforms for healthcare and financial services clients
- **AI/ML on AWS:** Built end-to-end MLOps pipelines using SageMaker — from feature engineering to model deployment and monitoring. Hands-on with SageMaker Pipelines, Model Registry, Feature Store, and Model Monitor
- **Generative AI:** Architected RAG-based solutions using Amazon Bedrock and SageMaker for document processing and intelligent Q&A systems
- **DevSecOps:** Implemented IaC with Terraform across multi-cloud environments, CI/CD with GitHub Actions, and security scanning (tfsec, Checkov, Trivy)
- **Compliance:** Deep experience with HIPAA, PCI-DSS, and SOC2 compliance on AWS — encryption, audit logging, access controls

I'm particularly excited about EXL because of its strength in analytics and BFSI domain — my SageMaker and GenAI experience directly maps to building AI-powered solutions for banking, insurance, and financial services clients."

Q33: Describe a complex architecture you designed. What were the trade-offs?

Answer:

"I designed a multi-cloud healthcare platform (MedCloud Global) spanning AWS, Azure, and GCP with strict HIPAA/PCI-DSS compliance:

Challenge: Process patient medical records (PHI) on Azure, run e-commerce payments (PCI) on AWS, and perform analytics/ML on GCP — all interconnected with zero-trust security.

Architecture decisions and trade-offs:

Decision	Trade-off
Multi-cloud (vs. single cloud)	Operational complexity ↑, but regulatory compliance met + vendor lock-in ↓
Istio mTLS mesh (vs. API gateway)	Higher infrastructure overhead, but end-to-end encryption + fine-grained auth policies
SageMaker on AWS (vs. Vertex AI on GCP)	Required cross-cloud data movement, but SageMaker's MLOps tooling (Pipelines, Model Monitor) was more mature
Aurora PostgreSQL (vs. DynamoDB)	Higher cost per query, but needed complex joins for order processing + ACID guarantees

Outcome: Successfully deployed 6 microservices across 3 clouds with automated CI/CD, passing compliance audits for HIPAA and PCI-DSS."

Q34: How do you handle a situation where the client wants a quick PoC but you know the architecture won't scale?

Answer:

"I use a '**PoC with Production Guardrails**' approach:

1. **Build the PoC fast** using managed services (Bedrock for GenAI, SageMaker JumpStart for ML) — deliver in 2-3 weeks
2. **Document the scaling gaps** clearly in an architecture decision record (ADR)
3. **Embed production patterns from day one:** Even in PoC, use VPC endpoints, KMS encryption, and CloudTrail logging
4. **Present a migration path:** Show the client the PoC architecture alongside the production target architecture with a clear migration timeline

This way, the client gets fast results, but there are no surprises when we productionize. I never deliver a PoC that can't be evolved — I'd rather spend an extra 2 days upfront on the foundation than 2 months refactoring later."

Q35: How do you stay current with rapidly evolving AI/cloud technologies?

Answer:

- **AWS re:Invent & re:Mars:** Attend virtually, review all AI/ML session recordings
 - **AWS Blogs:** Follow the Machine Learning and Architecture blogs weekly
 - **Hands-on:** Personal AWS sandbox account for experimenting with new services (Bedrock Agents, SageMaker HyperPod)
 - **Certifications:** AWS Solutions Architect Professional, AWS Machine Learning Specialty
 - **Community:** AWS Community Builder, contribute to open-source Terraform modules
 - **Internal:** Run monthly "AI Architecture Office Hours" with the engineering team to share learnings
-

Q36: How do you handle disagreements with development teams about architecture decisions?

Answer:

"I follow a data-driven approach:

1. **Listen first:** Understand the team's concerns — they often have valid operational context I lack
2. **Evidence-based discussion:** Create a comparison matrix (performance, cost, operational complexity, security) for both approaches
3. **PoC when needed:** If the disagreement is fundamental, I propose a time-boxed PoC (1-2 weeks) to test both approaches with real metrics

4. **ADR documentation:** Record the decision, rationale, and what we chose *not* to do (and why)
5. **Disagree and commit:** Once a decision is made, I fully commit — no passive resistance

Example: A team wanted to use Lambda for ML inference (cost optimization), I recommended SageMaker endpoints (latency). We PoC'd both — Lambda cold starts were 15s (unacceptable for real-time fraud), so the team agreed on SageMaker with a cost-optimization plan (Savings Plans + auto-scaling)."

Q37: How do you mentor junior architects and engineers?

Answer:

- **Architecture reviews:** Weekly 1-hour sessions where team members present their designs and I provide feedback
 - **Pair architecting:** Work side-by-side on complex designs (like VPC peering, SageMaker pipeline setup)
 - **Tech talks:** Monthly internal talks on new AWS services or architecture patterns
 - **Documentation:** Maintain an internal wiki with architecture decision records, design templates, and best practices
 - **Certification support:** Help team members prepare for AWS certifications (study groups, practice exams)
-

11. EXL-Specific Questions

Q38: Why EXL Services?

Answer:

"Three reasons:

1. **Analytics DNA:** EXL is ranked among the top analytics organizations globally, with 2000+ analytics professionals. My AI/ML architecture skills would be amplified by

this deep analytics capability — I can design the infrastructure and ML platforms, while EXL's data scientists build the models.

2. **BFSI Domain Depth:** EXL's focus on Banking, Financial Services, and Insurance aligns perfectly with my experience building compliant, secure AI solutions. The intersection of regulatory compliance and GenAI is where the hardest and most impactful problems are.
 3. **Digital Transformation Focus:** EXL's approach of 'looking deeper at the entire value chain' rather than just technology migrations resonates with me. I believe AI architecture should drive business outcomes, not just technology upgrades. Building an AI-powered claims processing system that reduces processing time from 5 days to 5 minutes — that's the outcome-oriented work I want to do."
-

Q39: What value would you bring as an AI Cloud Architect at EXL?

Answer:

"I bring three key capabilities:

1. **Hands-on SageMaker & GenAI expertise:** I've built production ML pipelines, not just designed them. I can set up SageMaker Pipelines, Model Monitor, Feature Store, and Bedrock-based RAG solutions from scratch — accelerating client delivery.
 2. **Enterprise architecture at scale:** Multi-account AWS landing zones, zero-trust networking, compliance automation (HIPAA, PCI-DSS) — I can architect solutions that pass regulatory audits from day one.
 3. **Cloud + AI bridge:** Many organizations have Cloud Architects who don't understand AI, and Data Scientists who don't understand cloud infrastructure. I bridge that gap — I can design the SageMaker training infrastructure *and* have meaningful conversations with data scientists about model architecture, hyperparameters, and deployment patterns."
-

Q40: How would you approach a new BFSI client wanting to adopt GenAI?

Answer:

GENAI ADOPTION ROADMAP FOR BFSI:

Phase 1: ASSESS (Week 1-2)

Identify top 5 use cases (claims processing, customer service, fraud, compliance, document processing)

Data readiness assessment (what data exists, quality, governance)

Security & compliance requirements (PII handling, model risk management)

Quick win identification (which use case delivers value fastest)

Phase 2: FOUNDATION (Week 3-6)

AWS Landing Zone setup (if not exists)

SageMaker Domain in VPC-only mode

Bedrock access with guardrails configured

Data pipeline (S3 → Glue → Feature Store)

Security controls (KMS, VPC endpoints, audit logging)

Phase 3: QUICK WIN POC (Week 7-10)

Build RAG-based policy Q&A chatbot (Bedrock + Knowledge Bases)

or Document processing pipeline (Textract + Bedrock)

Measure: time saved, accuracy, user satisfaction

Demo to stakeholders

Phase 4: SCALE (Month 3-6)

Custom model training (SageMaker) for domain-specific use cases

MLOps pipeline (SageMaker Pipelines + Model Registry)

Production deployment with monitoring

Expand to additional use cases

Quick Reference — Top SageMaker Services to Highlight

SageMaker Service	Your Experience	When to Mention
SageMaker Studio	Primary IDE for ML team, VPC-only mode	IDE/development environment questions
SageMaker Training	Distributed training, spot instances, Hugging Face integration	Training at scale questions

SageMaker Service	Your Experience	When to Mention
SageMaker Pipelines	End-to-end MLOps automation	CI/CD for ML questions
SageMaker Model Registry	Model versioning with approval workflows	Governance questions
SageMaker Feature Store	Online + offline stores for fraud detection	Feature engineering questions
SageMaker Model Monitor	Drift detection, bias monitoring	Production monitoring questions
SageMaker Clarify	Bias detection, SHAP explanations	Responsible AI questions
SageMaker Endpoint	Real-time, serverless, multi-model deployments	Deployment questions
SageMaker JumpStart	Foundation model evaluation and fine-tuning	GenAI / LLM questions
SageMaker Processing	Spark-based feature engineering jobs	Data processing questions

Interview Tips for EXL L1 Round

1. **Lead with architecture diagrams** — draw/describe system designs proactively
2. **Always connect to BFSI** — every answer should reference banking, insurance, or financial services
3. **Emphasize SageMaker hands-on** — mention specific instance types, API calls, and real metrics
4. **Show GenAI depth** — RAG, Bedrock Guardrails, prompt engineering — these are hot topics
5. **Talk cost optimization** — EXL clients care about ROI; mention Spot training, Savings Plans
6. **Security first** — BFSI = regulation; always mention encryption, audit logging, compliance

7. **Be outcome-oriented** — EXL's DNA is outcomes, not just technology. Frame answers as business impact

Prepared: June 2026 | **Candidate:** Pushparaj Naik | **Role:** AI Cloud Architect — EXL Services