

- Nextrun — Multi-Cloud & DevOps Architect
- L3 Techno-Managerial Round Interview Preparation
 - Why Nextrun Added an L3
 - What Makes L3 Different
 - Table of Contents
 - 1. Platform Strategy & Organizational Leadership
 - Q1: "You're hired as the Multi-Cloud & DevOps Architect. It's Day 1. What's your 30-60-90 day plan?"
 - Q2: "We currently use all three clouds — AWS (60%), Azure (25%), GCP (15%). Should we consolidate? How would you approach this decision?"
 - Q3: "How would you design a platform engineering organization from scratch for our 150-developer company?"
 - 2. Delivery Management & Execution
 - Q4: "You're tasked with building an Internal Developer Platform. How do you scope, plan, and execute this as a project?"
 - Q5: "A critical cloud migration project is 2 months behind schedule. You're brought in to rescue it. What do you do?"
 - 3. Team Building, Hiring & Talent Development
 - Q6: "How do you build a high-performing multi-cloud platform team? What do you look for in hires?"
 - Q7: "How do you handle performance issues in your team? Walk me through a specific approach."
 - 4. Stakeholder Management & Influence
 - Q8: "The VP of Engineering wants to standardize on Azure for everything because of an existing Microsoft EA. Your technical assessment says AWS is better for the K8s platform workloads. How do you handle this?"
 - Q9: "Engineering teams are complaining that your platform team is a bottleneck. 'We file a ticket and wait 2 weeks for a namespace.' How do you fix this?"
 - 5. FinOps, Budgets & Commercial Awareness
 - Q10: "Our cloud bill is \$2.5M/year and growing 30% annually. Leadership wants you to reduce it by 25% without impacting performance. Present your plan."
 - 6. Incident Management & Operational Maturity
 - Q11: "A multi-cloud outage occurs — AWS EKS is down in us-east-1 and Azure AKS in eastus simultaneously. Walk me through your response as the most senior technical leader."

- 7. Vendor Management & Technology Strategy
 - Q12: "How do you evaluate and select tools for the platform? Give me your framework for choosing between, say, ArgoCD vs. Flux, or Datadog vs. Grafana."
- 8. Scenario-Based Situational Judgement
 - Q13: "A developer pushes a configuration change directly to production, bypassing CI/CD. It works fine. But it set a precedent. What do you do?"
 - Q14: "A team member privately tells you they've discovered a security vulnerability in production — an S3 bucket with customer PII is publicly accessible. The bucket was created by a VP's direct report. What do you do?"
- 9. Self-Awareness, Growth & Culture Fit
 - Q15: "What's the biggest mistake you've made as an architect, and what did you learn?"
 - Q16: "Rate yourself on these dimensions and explain honestly."
- 10. Rapid-Fire Techno-Managerial Questions
- 11. Questions YOU Should Ask
- Pre-Interview Power Prep
 - 48 Hours Before
 - Day Of
- The L3 Meta-Framework

Nextron — Multi-Cloud & DevOps Architect

L3 Techno-Managerial Round Interview Preparation

Round: L3 — Techno-Managerial (Unplanned Additional Round) **Interviewer:** Engineering Director / VP Engineering / CTO **Format:** 60–90 min | Leadership + Technical Judgement + Organizational Design **Builds on:** [Part 1 — Q1–Q38](#) + [Part 2 — Q39–Q82](#) **What L3 tests:** Can you lead a multi-cloud platform organization, not just design one?

Why Nextron Added an L3

An unplanned round typically signals one of these:

Signal	What They're Checking
Strong L1+L2 performance	They liked you; now a senior leader wants to validate "fit"
Role is more senior than originally scoped	They're considering you for a broader mandate
Cultural/leadership validation	Can this person manage client relationships, team dynamics, vendor negotiations?
Competing candidates	Tiebreaker round — differentiate on leadership, not just tech

Your advantage: You've already proven technical depth (82 Q&As). L3 is about demonstrating that you can **own the cloud strategy for the organization, not just execute tasks.**

What Makes L3 Different

L1/L2 (Technical Deep-Dive)	L3 (Techno-Managerial)
"Design a CI/CD pipeline for EKS"	"3 teams refuse to adopt your CI/CD standard. What do you do?"
"Compare Istio vs Linkerd"	"The CTO wants to scrap the service mesh investment. How do you respond?"
"How does Karpenter bin-packing work?"	"Our AWS bill doubled last quarter. Present your plan to leadership."
Right technical answers	Right decisions under ambiguity and organizational pressure

Table of Contents

1. Platform Strategy & Organizational Leadership
 2. Delivery Management & Execution
 3. Team Building, Hiring & Talent Development
 4. Stakeholder Management & Influence
 5. FinOps, Budgets & Commercial Awareness
 6. Incident Management & Operational Maturity
 7. Vendor Management & Technology Strategy
 8. Scenario-Based Situational Judgement
 9. Self-Awareness, Growth & Culture Fit
 10. Rapid-Fire Techno-Managerial Questions
 11. Questions YOU Should Ask
-

1. Platform Strategy & Organizational Leadership

Q1: "You're hired as the Multi-Cloud & DevOps Architect. It's Day 1. What's your 30-60-90 day plan?"

Answer:

DAY 1-30: LISTEN, OBSERVE, UNDERSTAND

— Week 1: Meet every engineering lead (1:1s)

- "What's your biggest infrastructure pain point?"
- "How long does it take you to deploy to production?"
- "What wakes you up at night?" (on-call, reliability)
- "What would you change about our platform if you could?"

— Week 2: Technical assessment

- Review: existing Terraform code, CI/CD pipelines, K8s configurations
- Run: Prowler/Wiz scan – quantify security posture
- Analyze: CloudWatch/Azure Monitor – baseline performance metrics
- Review: last 10 post-mortems – pattern analysis
- Map: what runs where? AWS accounts, Azure subscriptions, GCP

projects

- Week 3: Cost & compliance audit
 - AWS Cost Explorer + Azure Cost Management – top 20 cost drivers
 - Tag compliance: what % of resources are properly tagged?
 - Compliance gaps: CIS benchmarks, SOC 2 controls
 - Unused resources: idle EC2s, unattached EBS, over-provisioned RDS
- Week 4: Current-state architecture document
 - As-is architecture diagram (what actually exists, not what docs say)
 - Pain point heatmap (severity × frequency)
 - Technical debt register (categorized: security, cost, reliability, DX)
- Present findings to leadership (no recommendations yet – just observations)

DAY 31–60: QUICK WINS + STRATEGY

- Quick wins (visible impact, low risk):
 - Fix top 5 security findings (open security groups, missing MFA)
 - Enable cost anomaly alerts (AWS CE Anomaly Detection)
 - Set up Terraform formatting + validation in CI (if missing)
 - Create shared Slack channel for platform updates
 - Publish "How to request infrastructure" guide (reduce ticket noise)
- Draft cloud strategy document:
 - Multi-cloud rationale: why we use AWS + Azure + GCP (or should we reduce?)
 - Target-state architecture (18-month vision)
 - Platform team charter: what we own, what teams own
 - Technology radar: adopt, trial, assess, hold
 - Investment priorities (ranked by business impact)
- Socialize: Present draft to leadership and senior engineers for feedback (Don't finalize in isolation – co-create with the organization)

DAY 61–90: EXECUTE FIRST INITIATIVE

- Pick ONE high-impact initiative from the strategy:
 - Option A: "Golden path" – self-service new-service setup
 - Option B: Unified CI/CD pipeline standard (reduce 15 pipeline patterns to 3)
 - Option C: FinOps program – 30% cost reduction with visible dashboard
 - Option D: Security hardening sprint – close top 20 critical findings
- Deliver it end-to-end:
 - Design → build → deploy → measure impact → present results
 - This establishes credibility: "the new architect delivers, not just presents"
- Establish ongoing rhythms:
 - Weekly platform office hours (engineers come with questions)
 - Monthly architecture review (cross-team alignment)
 - Quarterly technology radar update
 - Monthly cost review with engineering leadership

Q2: "We currently use all three clouds — AWS (60%), Azure (25%), GCP (15%). Should we consolidate? How would you approach this decision?"

Answer:

"This is a strategic decision, not a technical one. I'd never recommend consolidation without understanding the WHY behind the current distribution.

My analysis framework:

STEP 1: UNDERSTAND WHY WE'RE MULTI-CLOUD TODAY

- └ Intentional choice? ("Azure for M365 integration, AWS for compute, GCP for analytics")
- └ Organic growth? ("Each team picked their own cloud")
- └ M&A? ("Acquired company runs Azure")
- └ Vendor negotiation? ("Committed spend across providers for leverage")
- └ If organic → consolidation opportunity. If intentional → optimize, don't fight it.

STEP 2: QUANTIFY THE COST OF MULTI-CLOUD

- └ Skills cost: How many engineers are proficient in all 3? (Usually < 20%)
 - Cross-training cost or hiring premium for multi-cloud talent
- └ Tooling cost: Are we running 3x monitoring, 3x CI/CD, 3x IaC?
 - Or is it unified (Terraform, Grafana, GitHub Actions)?
- └ Egress cost: How much data moves between clouds?
 - \$0.08-0.12/GB adds up fast at scale
- └ Operational overhead: 3 different IAM models, networking, compliance processes
 - └ Total "multi-cloud tax" – I'd estimate this in \$/year

STEP 3: QUANTIFY THE VALUE OF MULTI-CLOUD

- └ Best-of-breed services:
 - └ GCP 15% → Is it BigQuery? That's genuinely best-in-class. Keep it.
 - └ Azure 25% → Is it M365/Entra ID? Essential. Keep it.
 - └ If it's just "someone picked Azure for a VM" → candidate for consolidation
- └ Risk mitigation: Does one cloud's outage take us fully down?
- └ Vendor leverage: Does multi-cloud give us pricing power?
- └ Regulatory: Do some workloads NEED to be on a specific cloud?

STEP 4: MY RECOMMENDATION (typical)

- └ Consolidate to "Primary + Specialist" model:
 - └ AWS (Primary, 80%): All compute, K8s, databases, CI/CD
 - └ Azure (Specialist, 15%): M365, Entra ID, Azure AD integration
 - └ GCP (Specialist, 5%): BigQuery for analytics (if genuinely needed)
 - └ Migrate the 15-20% of workloads that are on the "wrong" cloud

— Savings estimate:

- Reduced tooling: \$X/year (fewer monitoring, fewer pipelines)
- Team efficiency: ~20% faster delivery (one platform to learn)
- Consolidation leverage: Larger AWS Savings Plan = deeper discount
- Typical: 15–25% total cloud cost reduction through consolidation

— What I'd NEVER do:

- Force consolidation if there's a genuine best-of-breed reason.
"Move BigQuery to Redshift because AWS is primary" = wrong.
"Move a random Azure VM running nginx to AWS" = right.


Q3: "How would you design a platform engineering organization from scratch for our 150-developer company?"

Answer:

PLATFORM TEAM DESIGN (Team Topologies approach):

SIZE: 8–12 people (5–8% of total engineering headcount)

STRUCTURE:



INTERACTION MODEL:

Platform team is a PRODUCT team. Their customer is internal developers.

- Product mindset: track NPS, feature requests, adoption metrics
- Self-service > tickets: if devs file tickets for environments, we've failed

- Measure success: "What % of infrastructure requests are self-served?"
Month 1: maybe 20% → Month 6 target: 80%
- Office hours, not on-call for dev teams: we teach, we don't do it for you

WHAT THE PLATFORM TEAM OWNS vs. DOESN'T OWN:

Platform Team OWNS	Application Teams OWN
K8s cluster lifecycle	Application code
Terraform modules	Application-level Terraform
CI/CD pipeline templates	Pipeline configuration
Monitoring infrastructure	Application dashboards/alerts
Security baselines	Application security fixes
Network infrastructure	Service-to-service connectivity
Cost governance tools	Cost optimization of their app
Incident process	On-call for their services

HIRING ORDER (if building from zero):

1. Senior Cloud/Platform Engineer (IaC + K8s → foundational)
2. Senior DevOps/CI-CD Engineer (pipelines + developer experience)
3. SRE / Security Engineer (observability + compliance)
4. Platform Product Manager (often neglected – but essential for prioritization)

2. Delivery Management & Execution

Q4: "You're tasked with building an Internal Developer Platform. How do you scope, plan, and execute this as a project?"

Answer:

PHASE 1: DISCOVERY (Weeks 1–3)

- Developer surveys: "What takes too long?" "What's painful?"
- Timing studies:
 - Time to first deployment for a new engineer: ____ days
 - Time from commit to production: ____ hours
 - Time to provision a new environment: ____ days
 - Time to debug a production issue: ____ hours
 - These become our BEFORE metrics (measure improvement later)
- Common patterns across 150 developers:

- 80% use Python/Java/Node (focus golden paths here)
 - 60% deploy to EKS, 25% to Azure AKS, 15% to GCP GKE
 - 10+ different CI/CD patterns (Jenkins, GitHub Actions, GitLab)
 - No standardized monitoring (each team rolls their own)
- IDP feature prioritization (MoSCoW):
- Must: Self-service environment provisioning, standardized CI/CD
 - Should: Developer portal (Backstage), cost visibility per team
 - Could: Service catalog, API documentation hub
 - Won't (now): Advanced self-service databases, ML platforms

PHASE 2: FOUNDATION (Weeks 4–8)

- Standardize CI/CD: 3 pipeline templates (Python, Java, Node)
 - Using GitHub Actions reusable workflows
 - Security scanning built-in (Trivy, Gitleaks, Checkov)
- Self-service Terraform: Module catalog
 - "I need a new EKS namespace" → terraform apply (5 min)
 - "I need a new RDS database" → terraform apply (15 min)
 - All modules include encryption, monitoring, tagging by default
- Early adopters: 2–3 volunteer teams use it first
 - Their feedback shapes the next iteration

PHASE 3: DEVELOPER PORTAL (Weeks 9–16)

- Backstage deployment:
 - Service catalog: every service registered with owner, docs, SLOs
 - Software templates: "Create new service" wizard
 - TechDocs: architecture decisions, runbooks, API docs
 - Plugins: Kubernetes, ArgoCD, Cost, PagerDuty
- Measure adoption:
 - % of new services created via portal (target: 80% by month 4)
 - Developer satisfaction survey (target: NPS > 40)
 - Time savings: before vs. after metrics

PHASE 4: SCALE & OPTIMIZE (Weeks 17+)

- Rollout to all 150 developers (with training sessions)
- FinOps integration: per-team cost dashboards in Backstage
- SLO dashboard: every service shows its error budget burn rate
- Continuous improvement: monthly retro with developer representatives
- Success criteria:
 - New engineer to first production deploy: < 1 day (was: 2 weeks)
 - Commit to production: < 30 min (was: 2 days)
 - Self-service rate: > 80% (was: < 10%)
 - Developer NPS: > 50

Q5: "A critical cloud migration project is 2 months behind schedule. You're brought in to rescue it. What do you do?"

Answer:

FIRST 72 HOURS: DIAGNOSE (not fix)

DIAGNOSIS FRAMEWORK:

1. SCOPE ANALYSIS

- Compare original SOW/scope vs. what's actually being built
- If scope grew 40%, no wonder it's 2 months late
- Action: identify what can be DEFERRED to Phase 2 (not "cut" – reframed)
- Common finding: "We're trying to modernize AND migrate at the same time"
- Fix: Migrate first (rehost/replatform), modernize later (refactor)

2. BLOCKERS ANALYSIS

- Review Jira board: how many tickets are "blocked" or "waiting"?
- Common blockers in cloud migrations:
 - Waiting for security approvals (IAM roles, network rules)
 - Waiting for client decisions (architecture choices, vendor selection)
 - Environment provisioning delays (manual processes)
 - Data migration complexity underestimated
- Unblock immediately: What can I resolve TODAY?
- Set up: daily 15-min "blockers-only" standup

3. TEAM ASSESSMENT

- Is the team right-sized? (often under-staffed for cloud migration)
- Is the team right-skilled? (Java developers doing Terraform is slow)
- Is one person doing 40% of the work? (bus factor = 1 → key-person risk)
- Action: bring in specialist short-term if skill gap is structural

4. ESTIMATION REALITY CHECK

- Were the original estimates realistic? (often no)
- What's the honest forecast now? (use PERT: optimistic/likely/pessimistic)
- Options for stakeholders:
 - a) Reduce scope → deliver 70% on time
 - b) Extend timeline → deliver 100% in 2 extra months
 - c) Add team → deliver 100% on time but costs more (+ ramp-up risk)
 - d) Hybrid: scope down + small extension + targeted staff augmentation

PRESENT TO LEADERSHIP (Day 3–5):

"Here's what happened, here's why, here are 3 options with trade-offs. I recommend Option D because [specific reasoning]. I take accountability for the recovery plan from this point forward."

KEY PRINCIPLE:

Never surprise stakeholders. If a project is behind, acknowledge it early, present options, and let leadership choose. Surprises destroy trust.

3. Team Building, Hiring & Talent Development

Q6: "How do you build a high-performing multi-cloud platform team? What do you look for in hires?"

Answer:

HIRING PHILOSOPHY:

WHAT I LOOK FOR (in order of importance):

1. PROBLEM-SOLVING MINDSET

"Tell me about a production issue you debugged that took more than 2 hours.

Walk me through your thought process."

→ I'm evaluating: structured thinking, not memorized answers

2. T-SHAPED SKILLS

Deep in ONE area (Kubernetes, Terraform, AWS networking)

+ broad curiosity across the stack

Red flag: "I only do Terraform" → that's a tool operator, not an architect

3. COMMUNICATION

"Explain Kubernetes networking to a product manager in 2 minutes"

→ If they can't explain it simply, they don't understand it deeply enough

4. OWNERSHIP MENTALITY

"Tell me about something you built that you're proud of – not at work."

→ Side projects, open source, blog posts, homelab – shows genuine passion

5. LEARNING TRAJECTORY

"What have you learned in the last 6 months that changed how you work?"

→ Growth mindset > current skill level

WHAT I DON'T CARE ABOUT:

└─ Certifications alone (they validate knowledge, not ability)

└─ Years of experience (10 years of repeating Year 1 ≠ 10 years of growth)

└─ Specific tool experience (Terraform takes 2 weeks to learn; problem-solving takes years)

└─ Pedigree (FAANG experience is nice, not required)

INTERVIEW PROCESS I DESIGN:

└─ Round 1: Technical screening (45 min)

└─ System design: "Design the CI/CD platform for 20 teams" (open-ended)

- | | — Debugging: "Here's a Terraform plan that's creating something unexpected. Find the bug."
- | | — Evaluate: HOW they think, not WHAT they know
- | — Round 2: Hands-on pairing (60 min)
- | | — Real task from our backlog: "Add monitoring to this Terraform module"
- | | | — Pair program together – I'm looking for:
 - | | | | — Do they read documentation when stuck? (good)
 - | | | | — Do they ask clarifying questions? (good)
 - | | | | — Do they jump to coding without understanding the problem? (bad)
- | | | — This tells me more than 10 whiteboard questions
- | — Round 3: Culture + leadership (45 min)
- | | — "Describe a time you disagreed with your manager's technical decision"
- | | — "How do you handle a teammate who consistently doesn't pull their weight?"
- | | — "What would you do if you discovered a security vulnerability on a Friday at 5 PM?"

TEAM COMPOSITION I TARGET (8-person team):

- | — 2 Senior Cloud Engineers (AWS + Azure deep expertise)
- | — 2 Platform Engineers (Kubernetes + CI/CD + Developer Experience)
- | — 1 SRE (Observability + incident management + on-call design)
- | — 1 Security Engineer (Policy-as-code, compliance, DevSecOps)
- | — 1 Mid-level Engineer (growth potential, mentorship target)
- | — 1 Platform Engineering Manager / Tech Lead (me, or someone I hire)

Q7: "How do you handle performance issues in your team? Walk me through a specific approach."

Answer:

"The key insight is that underperformance almost always has a root cause that isn't laziness.

DIAGNOSIS FIRST (not action):

STEP 1: DATA, NOT FEELINGS

- | — Review: PR activity, ticket completion rate, code review participation
- | — Compare: to their own historical performance (not to the team's best performer)
- | — Ask their teammates (carefully): "How's the collaboration going?"
- | — Is this new behavior or consistent? (new = something changed in their life/role)

STEP 2: PRIVATE 1:1 CONVERSATION

"I've noticed your delivery has slowed down in the last few weeks. I'm not here to criticize – I want to understand what's going on and see how I can help."

COMMON ROOT CAUSES AND RESPONSES:

Cause A: WRONG ROLE/ASSIGNMENT

- Backend developer doing infrastructure work they hate
- Solution: Reassign to something that uses their strengths
- This is MY failure as a manager, not theirs

Cause B: SKILL GAP

- They've never worked with Terraform/K8s before this project
- Solution: Pair programming with a senior engineer for 2 weeks + online course budget + dedicated learning time (20% of sprint)
- Set clear milestones: "By week 4, you should be able to create a module independently"

Cause C: PERSONAL ISSUES

- Health, family, burnout – it happens to everyone
- Solution: "Take the time you need. Let's redistribute your work temporarily. No questions asked."
- Follow up in 2 weeks: "How are you doing? Ready to ramp back up?"

Cause D: GENUINE FIT ISSUE

- After 4–6 weeks of support, performance hasn't improved
- Solution: Transparent conversation
 - "I've provided pairing, training, and adjusted your workload. I'm still seeing the same challenges. Let's discuss whether this role is the right fit for your skills and interests."
- Help them transition (internally or externally) with dignity
- Timeline: never drag this out beyond 6 weeks after intervention starts

WHAT I NEVER DO:

- Discuss performance issues in front of the team
- Compare them to other team members ("Why can't you be like Ankit?")
- Assume the worst without investigating
- Wait for annual review to address a problem that's visible now

4. Stakeholder Management & Influence

Q8: "The VP of Engineering wants to standardize on Azure for everything because of an existing Microsoft EA. Your technical

assessment says AWS is better for the K8s platform workloads. How do you handle this?"

Answer:

"This is a classic 'politics vs. technology' situation. The answer is almost never pure technology.

STEP 1: UNDERSTAND THE FULL PICTURE

- What does the Microsoft EA cover? (Compute credits? M365? Azure AD?)
- What's the discount? (Often 20–40% on Azure compute – that's real money)
- What's the VP's relationship with Microsoft? (Strategic partnership? Personal?)
- What are the switching costs? (Existing Azure skills, tools, integrations)
- Key question: "Is this a strong preference or a hard requirement?"

STEP 2: OBJECTIVE COMPARISON (for K8s specifically)

Criteria	EKS (AWS)	AKS (Azure)
Managed K8s maturity	★★★★☆	★★★★★ (free CP)
Node autoscaling	Karpenter (best)	Cluster autoscaler
Networking	VPC CNI	Azure CNI Overlay
Service mesh support	Istio (mature)	Istio (mature)
Graviton ARM instances	✅ (20% cheaper)	❌ (limited ARM)
Spot instance diversity	✅ (hundreds)	✅ (fewer types)
Team's existing skills	Strong	Moderate
EA discount	None	30%
Identity integration	Separate SS0	Native Entra ID
Cost (with EA)	\$X	\$X – 30%

STEP 3: PRESENT OPTIONS (not "my recommendation")

Option A: "All-in Azure AKS"

Pro: EA discount, single cloud, Entra ID native

Con: Give up Karpenter, Graviton, some AWS-specific services

Risk: Medium (AKS is mature enough for most workloads)

Option B: "Primary AWS + Azure for identity/M365" (my recommendation)

Pro: Best-of-breed K8s platform, maintain team's AWS skills

Con: Multi-cloud operational overhead, doesn't maximize EA

Risk: Low (proven pattern at scale)

Option C: "Start Azure AKS, evaluate in 6 months"

Pro: Respects VP's preference, limits downside with time-boxed evaluation

Con: Team ramp-up cost, potential migration if we switch back

Risk: Medium

"I recommend Option B, but Option C is a viable compromise that lets us validate AKS with real workloads before committing fully."

STEP 4: ACCEPT THE DECISION GRACEFULLY

If the VP chooses Azure: "Understood. Here's my plan to make AKS successful. I want to ensure we:

1. Invest in Azure training for the team (2-week intensive)
2. Use Terraform + Kubernetes (not Azure-specific tools) for portability
3. Set a 6-month checkpoint to evaluate AKS performance vs. baseline
4. Document this as an ADR with the rationale"

KEY PRINCIPLE:

"My job is to present data and options, not to win arguments.
A good architect makes ANY reasonable choice successful."

Q9: "Engineering teams are complaining that your platform team is a bottleneck. 'We file a ticket and wait 2 weeks for a namespace.' How do you fix this?"

Answer:

FIRST: ACKNOWLEDGE THE PROBLEM (don't get defensive)

"If teams are waiting 2 weeks for a namespace, we've failed.

A platform team that's a bottleneck is worse than no platform team."

ROOT CAUSE ANALYSIS:

- Why does a namespace take 2 weeks?
 - Manual process? → Automate with self-service Terraform module
 - Approval bottleneck? → Who approves? Why? Can we pre-approve for standard requests?
 - Security review? → Embed security defaults so review is unnecessary
 - Capacity issue? → Platform team has 50 tickets and 3 engineers
 - Knowledge gap? → Teams don't know how to do it themselves
- Categorize ALL platform team requests:
 - Category A (90%): Standard, repeatable (new namespace, new DB, new pipeline)
 - These MUST be self-service. Zero ticket, zero wait.
 - Category B (8%): Semi-custom (cross-account networking, special IAM)
 - Template + light review (< 2 business days SLA)
 - Category C (2%): Truly custom (new cloud region, compliance-specific)
 - Scheduled work (sprint planning, treated as a project)

SOLUTION: SELF-SERVICE PLATFORM

- For Category A (90% of requests):
 - Terraform module: "I need a namespace"
 - Developer runs: terraform apply → namespace + RBAC + NetworkPolicy + ResourceQuota

- Time: 5 minutes, zero human approval needed
- └─ Terraform module: "I need a database"
 - Developer runs: terraform apply → RDS/Aurora + secrets + monitoring
- Time: 15 minutes, zero human approval needed
- └─ Backstage template: "I need a new service"
 - Developer clicks wizard → repo + pipeline + K8s manifests + monitoring
 - Time: 10 minutes, zero human approval needed
- └─ For safety: Pre-approved guardrails (not gatekeeping)
 - └─ OPA/Kyverno policies prevent dangerous configurations automatically
 - └─ Cost limits: namespace can't exceed X CPU/memory (ResourceQuota)
 - └─ Security: all resources get encryption, logging, tagging by default
 - └─ If the guardrail blocks you, submit a Category B request with justification
- └─ MEASURE AND REPORT:
 - └─ Track: "% of requests self-served" (target: 90%)
 - └─ Track: "median time-to-resolution for Category B" (target: < 2 days)
 - └─ Monthly report: "Platform team handled X requests; Y% self-service"
 - └─ Publish SLAs: Category A = instant, Category B = 2 days, Category C = sprint planning

5. FinOps, Budgets & Commercial Awareness

Q10: "Our cloud bill is \$2.5M/year and growing 30% annually. Leadership wants you to reduce it by 25% without impacting performance. Present your plan."

Answer:

FINOPS ENGAGEMENT – 90-DAY PLAN:

WEEK 1-2: VISIBILITY

- └─ Enable AWS Cost Explorer hourly granularity
- └─ Enable Azure Cost Management + Cost Analysis
- └─ Tag audit: what % of spend is untagged? (typically 20-40%)

- Untagged = unaccountable. Fix this first.
- Install Kubecost on all K8s clusters (per-namespace cost allocation)
- Infracost in CI/CD (cost impact on every Terraform PR)
- TOP 10 cost drivers analysis (usually 80/20 rule applies)

TYPICAL FINDINGS AND SAVINGS:

Optimization	Savings	Implementation
Right-sizing (oversized VMs)	15-20%	AWS Compute Optimizer recs
Spot/preemptible for K8s	10-15%	Karpenter Spot
diversification Savings Plans (1-yr compute)	8-12%	Commit on baseline usage
Dev/staging auto-shutdown	5-8%	Lambda + EventBridge 7pm-7am
Unused resources cleanup	5-10%	Idle EC2s, unattached EBS
S3 lifecycle policies	2-5%	Intelligent-Tiering + Glacier
NAT Gateway optimization	2-4%	VPC Endpoints for AWS APIs
Right-sized RDS instances	3-5%	Scale down dev/staging DBs
TOTAL	30-45%	Exceeds the 25% target

AT \$2.5M/YEAR:

- 25% savings target = \$625K/year
- Achievable timeline: 50% of savings in 30 days, 100% in 90 days
- No performance impact if done correctly

GOVERNANCE (ongoing):

- Monthly cost review meeting (15 min, automated dashboard)
- Per-team cost dashboards (showback → eventual chargeback)
- Budget alerts at 80% and 100% of monthly allocation
- Anomaly detection: alert if daily spend spikes > 20%
- Quarterly Savings Plan optimization review
- FinOps champion per team (not platform team's sole responsibility)

PRESENTATION TO LEADERSHIP:

"Here's our \$2.5M cloud spend breakdown. I've identified \$750K in savings across 8 categories. I need 2 weeks for quick wins (\$300K) and 90 days for the full program (\$750K). Zero performance impact. The catch: ongoing governance is required. Without it, costs drift back within 6 months."

6. Incident Management & Operational Maturity

Q11: "A multi-cloud outage occurs — AWS EKS is down in us-east-1 and Azure AKS in eastus simultaneously. Walk me through your response as the most senior technical leader."

Answer:

THIS IS A P0 – LEADERSHIP RESPONSE, NOT JUST ENGINEERING:

MINUTE 0–5: MOBILIZE

- Acknowledge PagerDuty alert within 2 minutes
- Open war room (Zoom bridge + Slack #incident-YYYY-MM-DD)
- Declare myself as Incident Commander (IC)
 - IC DOES NOT DEBUG. IC coordinates.
 - Assign: "Raj, you own AWS diagnosis. Priya, you own Azure diagnosis."
 - Assign: Communications lead ("Amit, update Statuspage every 15 min")
- Quick blast radius assessment: which services? which customers? revenue impact?

MINUTE 5–15: PARALLEL TRIAGE

- AWS Team asks:
 - "Is this a regional AWS outage?" → Check health.amazonaws.com
 - "Recent deployment?" → Check ArgoCD + GitHub Actions history
 - "EKS control plane or data plane?" → `kubectl get nodes` (timeout = CP issue)
 - "Node health?" → CloudWatch → EC2 status checks
- Azure Team asks:
 - "Is this Azure-wide?" → Check status.azure.com
 - "AKS provisioning state?" → `az aks show`
 - "Node pool status?" → `kubectl get nodes`
 - "Networking?" → NSG flow logs, VPN connectivity to AWS
- CROSS-CLOUD: Is it our cross-cloud VPN link?
 - If both clouds lose connectivity to each other simultaneously → VPN issue
 - Check: IPsec tunnel status on both sides
- DECISION TREE:
 - If cloud provider outage → Failover to DR (or wait if DR is also affected)
 - If our deployment broke it → Rollback (ArgoCD: `argocd app rollback`)

- └─ If VPN issue → Route traffic to single-cloud mode
- └─ If unclear → Scale the war room (bring in more engineers)

MINUTE 15-60: MITIGATE

- └─ ALWAYS prefer mitigation over root cause analysis during active incident
"Stop the bleeding first. Understand the wound later."
- └─ Mitigation options (try in order):
 1. Rollback last deployment (if deployed in last 2 hours)
 2. Scale up healthy components (add nodes, increase replicas)
 3. Failover to DR region/cloud
 4. Feature-flag disable non-critical features (reduce load on healthy components)
 5. Traffic shed: return 503 for non-critical endpoints to protect critical ones
- └─ Communicate every 15 minutes:
 - Internal: Slack #incident channel
 - External: Statuspage / customer communication lead
 - Executive: direct message to VP/CTO with 2-sentence update

POST-INCIDENT (within 48 hours):

- └─ Blameless post-mortem (REQUIRED for P0)
 - └─ Timeline (factual, minute-by-minute)
 - └─ Root cause (5-whys)
 - └─ What went well (acknowledge good response)
 - └─ What could be improved
 - └─ Action items with OWNERS and DEADLINES
- └─ Systemic improvements:
 - └─ What monitoring gap existed? (should have alerted earlier)
 - └─ What runbook was missing? (create it now)
 - └─ What automation would have reduced MTTR?
 - └─ Do we need to reconsider our DR strategy?
- └─ Share learnings:
 - Present post-mortem to full engineering org (transparency builds trust)

7. Vendor Management & Technology Strategy

Q12: "How do you evaluate and select tools for the platform? Give me your framework for

choosing between, say, ArgoCD vs. Flux, or Datadog vs. Grafana."

Answer:

TOOL EVALUATION FRAMEWORK (5 dimensions):

1. TECHNICAL FIT (40% weight)

- └ Does it solve the ACTUAL problem? (not "cool technology looking for a problem")
- └ Integration with existing stack (K8s, Terraform, CI/CD, monitoring)
- └ Performance at our scale (PoC, not vendor promises)
- └ Feature completeness for OUR use case (not all features matter)
- └ Extensibility: can we customize/extend when we hit limits?

2. OPERATIONAL COST (25% weight)

- └ TCO over 3 years: license + infrastructure + team time
 - └ Datadog: \$400K/yr license + \$0 infra + 2 hours/week ops
 - └ Grafana OSS: \$0 license + \$50K/yr infra + 8 hours/week ops
 - └ Real comparison: Datadog \$1.2M vs Grafana \$250K over 3 years
 - └ But: Grafana requires a skilled SRE to maintain; Datadog doesn't
- └ Vendor lock-in risk: what happens if we need to switch?
 - └ OTel-based instrumentation = low lock-in (works with either)
 - └ Proprietary agents = high lock-in (expensive to migrate)
- └ Scaling costs: what does it cost at 10x our current scale?

3. TEAM CAPABILITY (20% weight)

- └ Current team skills: do we have people who know this?
- └ Hiring market: can we find engineers for this tool?
- └ Learning curve: how long to be productive? (weeks? months?)
- └ Community: Stack Overflow answers, GitHub issues, meetups

4. ECOSYSTEM & MOMENTUM (10% weight)

- └ Community size and growth rate (GitHub stars, contributors, releases)
- └ Enterprise adoption: are companies at our scale using it?
- └ Vendor stability: is the company profitable? (OSS: is the project active?)
- └ CNCF graduation status (for cloud-native tools)

5. COMPLIANCE & SECURITY (5% weight)

- └ SOC 2 Type II certification (for SaaS vendors)
- └ Data residency: where does data go? (for regulated industries)
- └ Audit trail: can we track who did what?
- └ BAA availability (for healthcare/HIPAA)

SPECIFIC EXAMPLE: ArgoCD vs. Flux

Criteria	ArgoCD	Flux
UI dashboard	✓ Rich UI (visual sync status)	✗ No built-in UI
Multi-cluster	✓ Native (ApplicationSet)	✓ (but more complex)
Helm support	✓ Native	✓ Native
RBAC	✓ Fine-grained project RBAC	Basic

Community	★★★★★ (15K GitHub stars)	★★★★★☆ (6K stars)
Resource usage	Medium (UI + API server)	Light (no UI overhead)
My recommendation	**ArgoCD for teams > 5 services**	**Flux for GitOps purists**

DECISION PROCESS:

1. Technical evaluation doc (1 page, comparison matrix)
2. PoC: 1-week hands-on with top 2 candidates on a real workload
3. Decision meeting: present PoC results to team (not just my opinion)
4. ADR: document the decision with rationale and review date
5. Pilot: 1-month trial with one team before org-wide rollout

8. Scenario-Based Situational Judgement

Q13: "A developer pushes a configuration change directly to production, bypassing CI/CD. It works fine. But it set a precedent. What do you do?"

Answer:

"The fact that it worked is actually the dangerous part. If it had failed, the lesson would be self-evident. Success creates a precedent: 'See? The pipeline is just bureaucracy.'"

STEP 1: UNDERSTAND THE CONTEXT (private conversation, same day)

"Hey [developer], I noticed the prod config change was pushed directly. I want to understand – was the pipeline broken? Was there an urgency I should know about?"

LIKELY RESPONSES AND MY REACTION:

If "Pipeline was broken/slow":

└─ "That's a platform team failure. Let me fix the pipeline.
└─ In the meantime, here's how to fast-track: [process].
└─ But direct push isn't the workaround, even if the pipeline is slow."
└─ Action: Fix the pipeline issue FIRST (remove their justification)

If "It was urgent, customer was impacted":

- └─ "I understand urgency. Let's create an emergency path:
 - Emergency hotfix pipeline: runs security scan + deploys in 5 min
 - Still automated, still auditable, just faster
 - Requires post-deployment review within 24 hours"
- └─ Action: Build the emergency path so this never happens again

If "I didn't think it was a big deal":

- └─ "I appreciate the honesty. The change worked, but here's the risk:
 - No audit trail (compliance failure)
 - No rollback capability (if it broke something subtle)
 - Other developers will follow this precedent
 - Production changes without security scanning = liability
- └─ This isn't about trust in you – it's about trust in the process."
- └─ Action: Reinforce, don't punish. First offense = conversation, not consequence.

STEP 2: SYSTEMIC FIX (prevent, don't just scold)

- └─ Technical controls:
 - └─ ArgoCD: disable manual sync in production (only auto-sync from Git)
 - └─ K8s RBAC: remove direct kubectl apply permissions for prod namespace
 - └─ AWS: SCPs that prevent direct resource creation outside Terraform
 - └─ Git branch protection: require PR + approval for main branch
- └─ Create the "fast path" they need:
 - └─ Hotfix pipeline: abbreviated but still scanned (5 min end-to-end)
 - └─ Feature flags: change behavior without deployment
- (LaunchDarkly/ConfigMap)
 - └─ Emergency runbook: "If you need to bypass pipeline, here's the process"
 - (Requires: 2 approvals + Slack notification + post-deploy review)
- └─ Communication to the team:
 - "I've tightened production access controls. This isn't because I don't trust you – it's because production reliability requires a process that protects all of us. I've also created a hotfix pipeline for urgent changes.
 - Questions?"

KEY PRINCIPLE:

"Don't make the wrong thing easy. Make the right thing easier."

Q14: "A team member privately tells you they've discovered a security vulnerability in production — an S3 bucket with customer PII is publicly accessible. The bucket was created by a VP's direct report. What do you do?"

Answer:

THIS IS NOT A POLITICAL SITUATION. THIS IS A SECURITY INCIDENT.

MINUTE 0-15: VERIFY AND CONTAIN

- Verify the finding myself (aws s3api get-bucket-policy + get-public-access-block)
- If confirmed public: IMMEDIATELY block public access
 - aws s3api put-public-access-block --bucket [bucket] \
 - public-access-block-configuration \
 - BlockPublicAcls=true,IgnorePublicAcls=true,\
 - BlockPublicPolicy=true,RestrictPublicBuckets=true
- Check: Has the bucket been accessed by external IPs?
 - S3 access logs or CloudTrail data events
 - If external access found → this is a data BREACH, not just a vulnerability
- Document everything with timestamps

MINUTE 15-60: ESCALATE PROPERLY

- Notify: Security team / CISO immediately
 - "I've identified and contained a public S3 bucket containing PII. Access has been blocked. I need help assessing whether external access occurred."
- Notify: My manager / VP Engineering
 - "Flagging a security incident. Contained. Escalating per incident process."
- DO NOT: Go to the VP whose report created it first
 - The security team should manage stakeholder communication
 - If I go to the VP directly, they might pressure me to minimize or "handle it quietly" – that's not an option with PII
- The team member who reported it: "Thank you. I've escalated it. You did the right thing. I'll keep you updated."

POST-CONTAINMENT:

- Security team assesses: data breach or vulnerability?
 - If BREACH (external access confirmed):
 - Legal team involved (data breach notification laws: GDPR 72 hours)
 - Affected customers notified per regulatory requirements
 - Incident report filed
 - If VULNERABILITY (no external access):
 - Document in risk register
 - No customer notification required
 - Still treated as a serious finding
- Root cause: WHY was this bucket public?
 - No SCP preventing public S3? → Implement SCP immediately
 - No Terraform review caught it? → Add Checkov to CI/CD
 - Created manually (not via IaC)? → Enforce IaC-only provisioning
 - Person didn't know better? → Training, not blame
- Systemic fix:
 - AWS Config rule: auto-remediate any public S3 bucket (tag-based exceptions only)
 - SCP: deny s3:PutBucketPolicy with public Principal
 - Weekly scan: Prowler/Wiz detects public resources

└ Training: "Cloud security basics" mandatory for all engineers

REGARDING THE VP'S REPORT:

└ This is handled through normal incident response – not political channels

└ The person who created the bucket needs training, not punishment (unless they deliberately circumvented security controls)

└ If the VP pressures me to minimize: "I understand your concern, but PII exposure has regulatory implications. The security team is handling this by the book. That protects everyone, including your team."

└ Never cover up. Never minimize. The cover-up is always worse than the incident.

9. Self-Awareness, Growth & Culture Fit

Q15: "What's the biggest mistake you've made as an architect, and what did you learn?"

Answer:

"I once designed a multi-cloud platform where EVERY service was abstracted to be cloud-agnostic — Kubernetes everywhere, no managed services, everything containerized. The goal was 'run anywhere.'

The result? We reinvented AWS services poorly. We ran PostgreSQL on K8s instead of using Aurora (spent 20% of ops time managing database failovers that Aurora handles automatically). We self-hosted Kafka instead of using MSK (required a dedicated Kafka admin). We built our own secrets management instead of using Secrets Manager.

Total cost of 'cloud-agnostic': 3 extra engineers just to manage what managed services do better. \$200K+/year in unnecessary operational overhead. Zero benefit — we never actually moved between clouds.

What I learned:

1. CLOUD-AGNOSTIC $\neq$ CLOUD-PORTABLE
Using Kubernetes and Terraform gives you enough portability.
You don't need to avoid EVERY managed service.
2. ABSTRACTION HAS A COST
Every abstraction layer you add is a layer your team must debug, maintain, and upgrade. Ask: "What does this abstraction buy us that we will ACTUALLY use?"
3. THE RIGHT QUESTION IS: 'WHAT'S OUR ACTUAL LOCK-IN RISK?'
Database engine? Low risk (Aurora PostgreSQL $\rightarrow$ Azure PostgreSQL = compatible)
S3 API? Low risk (MinIO, Azure Blob with S3-compatible API)
Lambda-specific code? Higher risk – but often worth it for the productivity
4. USE MANAGED SERVICES BY DEFAULT. SELF-HOST ONLY WHEN JUSTIFIED.
Justified: compliance requirement, feature gap, cost (at massive scale)
Not justified: "what if we need to move" (you probably won't)

Q16: "Rate yourself on these dimensions and explain honestly."

Dimension	Rating	Honest Explanation
AWS depth	8.5/10	"Deep hands-on: networking, EKS, IAM, security services, cost optimization. Not a 10 — I don't claim expertise in niche services like AppSync or Braket."
Azure	7.5/10	"Strong: AKS, VNet, Entra ID, Key Vault. Gap: less experience with Azure-specific data services (Synapse, Cosmos DB at scale)."
GCP	6.5/10	"Competent: GKE, BigQuery, IAM. Haven't run production GCP infrastructure at scale — mostly used for analytics workloads."
Kubernetes	9/10	"This is my strongest area. Production at scale, CKA-level knowledge, Karpenter, Istio, GitOps, security hardening. The 1 point gap: I haven't written a K8s operator from scratch."
Terraform/IaC	9/10	"Module design for 50+ teams, state management at scale, OPA policy integration. Gap: haven't used

Dimension	Rating	Honest Explanation
		Pulumi/CDK deeply."
People leadership	7.5/10	"Led teams of 8-12, mentored architects, handled underperformers. Growth area: I sometimes jump into execution when I should be delegating."
Stakeholder management	8/10	"Can present to VPs and CTOs, handle difficult conversations. Growth area: being more concise in executive communication — I tend to over-explain technical details."
Architecture strategy	8.5/10	"Strong: multi-cloud strategy, platform design, technology evaluation. I think in trade-offs, not absolutes. Gap: limited experience with mainframe modernization."

10. Rapid-Fire Techno-Managerial Questions

Question	Strong Answer
"Your best engineer gets a 50% higher offer. What do you do?"	"Have an honest conversation. Understand the full picture — money, growth, challenge. Counter with what I can: promotion, learning budget, interesting projects, conference talks. If they leave, part well — and fix the systemic issue (are we paying market rate?)."
"How do you keep up with technology changes?"	"Curated, not exhaustive. Weekly: CNCF newsletters, The New Stack, Hacker News. Monthly: hands-on lab with one new tool. Quarterly: conference or workshop. I don't try to know everything — I try to know what matters for the next 12 months."
"How do you handle a decision you disagree with from above?"	"Disagree openly in private ('Here's my concern, here's the data'). If the decision stands, commit fully ('I'll make this succeed'). Never undermine a decision you've lost publicly."

Question	Strong Answer
"What's the most underrated cloud service?"	"VPC Endpoints / PrivateLink. Everyone focuses on compute and storage. But removing NAT Gateway data processing charges and eliminating internet exposure for internal AWS traffic saves 10-15% on networking costs alone."
"How do you measure platform team success?"	"Developer NPS, self-service rate, time-to-first-deploy for new engineers, DORA metrics, and one I always add: 'How many tickets did we NOT get this month?' (fewer tickets = better platform)"
"What's your approach to technical debt?"	"Treat it like financial debt — some is healthy (strategic shortcuts), most is dangerous (neglected maintenance). I dedicate 20% of every sprint to debt reduction. Non-negotiable. Track it visibly on a debt register, not hidden in the backlog."
"A team wants to use a new tool you've never heard of. What do you do?"	"Ask them to present: problem it solves, alternatives considered, TCO, team support plan. If it passes the evaluation framework, run a time-boxed PoC. I don't block innovation — I channel it."
"How do you balance security with developer velocity?"	"Make the secure path the easy path. If security requires 5 extra steps, developers will skip them. If security is built into the golden path pipeline, developers don't even notice it's there."
"What's the biggest waste of money you've seen in cloud?"	"Running dev and staging environments 24/7 when they're used 8 hours/day, 5 days/week. That's 76% waste. A simple Lambda + EventBridge auto-shutdown saves \$500K/year at scale."
"How do you handle giving negative feedback?"	"Directly, privately, and specifically. Not 'your work is below standard.' Instead: 'The Terraform module you wrote doesn't have variable validation. Here's an example of what good looks like. Let me pair with you on fixing it.'"

11. Questions YOU Should Ask

About the Role:

1. *"What does success look like for this role at 6 months and 12 months?"*
2. *"Is this role primarily hands-on technical, or does it include people management and team building?"*
3. *"How many engineers would I be working with directly, and what's the current team maturity level?"*

About the Organization:

4. *"What's the biggest infrastructure or platform challenge the team is facing today?"*
5. *"How are technology decisions currently made — is there an architecture review board, or is it distributed?"*
6. *"What's the current DevOps maturity level? Where are you on the DORA metrics spectrum?"*

About Multi-Cloud Specifically:

7. *"Is the multi-cloud strategy intentional and strategic, or has it grown organically? Are there plans to consolidate?"*
8. *"What's the current cloud spend, and is FinOps an active practice or something you want to establish?"*

About Growth:

9. *"How does Nextron support continued learning — conference budgets, certification support, experimentation time?"*
10. *"What does the career path look like beyond this role — VP Engineering? Distinguished Engineer? CTO track?"*

Pre-Interview Power Prep

48 Hours Before

- ☐ Re-skim Part 1 (Q1–Q38) and Part 2 (Q39–Q82) — don't deep-read, just refresh
- ☐ Practice 3 scenarios out loud (3 min each, timed):
 - Scenario: Team complaining platform is a bottleneck (Q9)
 - Scenario: Cloud bill doubled (Q10)
 - Scenario: Developer bypassed CI/CD (Q13)
- ☐ Prepare 2 STAR stories:
 - Story 1: Turning around a struggling project/team
 - Story 2: Making a tough call under pressure (security, architecture, people)
- ☐ Research the interviewer on LinkedIn — find common ground

Day Of

- ☐ **Opening:** "I'm excited to discuss how my multi-cloud platform experience can contribute to Nextron's engineering culture and technical direction."
- ☐ **When unsure:** "I haven't faced that exact situation, but here's how I'd approach it based on my experience with similar challenges..."
- ☐ **Show both sides:** For every answer, show technical depth AND leadership maturity. "Here's what I'd build (technical) and here's how I'd get the organization to adopt it (leadership)."
- ☐ **Close strong:** "This conversation has confirmed my excitement about this role. The multi-cloud platform challenge, combined with the opportunity to build and lead a team, is exactly the kind of impact I want to make. I'm ready to hit the ground running."

The L3 Meta-Framework

L1/L2 proved: You can design and implement multi-cloud platforms at scale. **L3 proves:** You can lead the PEOPLE, manage the STAKEHOLDERS, and make the DECISIONS that make those platforms successful in a real organization.

The answer they want to hear (through your answers, not explicitly):

"I don't just build platforms. I build the teams that build platforms. I don't just make technical decisions. I make decisions that balance technology, people, cost, and business outcomes. And when things go wrong — because they will — I own the problem, not just the architecture."

*Prepared for: Pushparaj Naik | Role: Multi-Cloud & DevOps Architect — Nextron |
Round: L3 Techno-Managerial Builds on: Part 1 (Q1–Q38) + Part 2 (Q39–Q82) = Total
prep: 98+ Q&As across 3 rounds Prepared: June 2026*